

FFORMA: Feature-based FORecast-model Averaging

Abstract

To produce good forecasts of time series, we must first solve the problem of which model to use. Moreover, forecasting accuracy can even be improved by combining different models. We present an automated method for combining forecasting models that outperforms any individual method in a set of the most popular ones in the time series forecasting literature, achieving second position in the M4 Competition. Our approach works in two phases. In the first, we use a dataset of time series to train a meta-model to assign probabilities to the forecasting methods with the goal of minimizing the average forecasting error. In the second phase, we forecast new series by assigning probabilities to methods using our previously trained meta-model, and combining their individual forecasts using a weighted average. The inputs to this model are features extracted from the series.

Keywords: FFORMA (Feature-based FORecast-model Averaging), FFORMS (Feature-based FORecast-model Selection), Time series features, Forecast combination, XGBoost, M4 Competition

1 Introduction

There are essentially two general approaches to forecast a time series: i) use of single model and ii) combination forecast or forecast model averaging. There is a growing consensus that the combination forecast is a way of improving forecast accuracy. Empirical work often provides evidence that the combination of forecasts resulting from several candidate forecasting models are often superior to their individual forecasts. Combining forecasts across different models is considered as a way of reducing the risk of selecting an inappropriate model. However, the main challenge in forecast combination is the selection of appropriate set of weights.

Bates & Granger (1969), was the first to forward the idea of “the combination of forecasts”. Since then, it has been extended in various ways and many approaches have been proposed to derive weights for forecast combination (Timmermann (2006)). Simple averages(ref_clemen1989), regression-based approaches, ... to name a few.

Recently, a few researchers have explored the use of time series features in selecting the most appropriate forecasting method. However, the use of time series features to derive weights for forecast combination has been rarely addressed in the literature. In this paper we propose a general framework to obtain weights for forecast combination based on features computed from the time series. We call this framework FFORMA (Feature-based FOfRecast Model Averaging). The proposed FFORMA framework has been used over the course of M4 competition and placed in second in forecast accuracy in both point forecasts and prediction intervals. This approach has been influenced by the work of Talagala, Hyndman & Athanasopoulos (2018) and is related to the previous work by ref_prudencio in which they use machine learning techniques to define weights for the linear combination of forecasts.

The rest of the paper is organized as follows: In Section 2 we describe the FFORMA framework in a general sense, without specifying the set of features, forecasting methods and learning model implementation. This section contains the main contribution: posing the learning problem in a way that includes the information about the errors produced by each forecasting method and combines it with the features extracted from the series. Section 3 gives the details of our implementation of FFORMA which achieved second place in the M4 Competition. The required preprocessing steps, the sets of features and forecast methods, as well as the specific implementation of the meta-learning model. *The method to produce the prediction intervals for the M4 competition is also explained in this section.* We show empirical evidence on the good performance of the approach in Section 4 by quantifying the difference between our proposal and a traditional classifier approach using the same features and underlying learning implementation.

2 Methodology

2.1 Overview / Intuition

The objective of our meta-learning approach is to combine a set of individual forecasting methods (e.g. ARIMA, exponential smoothing) to produce more accurate forecasts. This combination is a weighted average of the forecasts of individual methods. Our framework requires a dataset of time series, called the *reference set*, from which we learn to produce these weights for any given time series.

We start from the reference set and a fixed set of forecasting methods, called the *pool of methods*. The forecasting errors of each method in the pool are calculated for each time series in the reference set, using the true future values of the series if available or by dividing the series into

a training period and test period by applying temporal holdout.

The inputs for training the meta-learning model are:

- features (e.g. length, autocorrelation coefficients,...) extracted from the training period of the series.
- The errors produced by the forecasting methods.

The model will learn to produce weights for each method in the pool as a function of the features of the series. Once the meta-learning model is trained, the inductive step consist of assuming that a new series we want to forecast comes from a *generating process* similar to the one that generated the reference set.

One common meta-learning approach is to learn to select the best method in the pool for each series, the one that produces the least forecasting error. This approach transforms the problem into a traditional classification one by setting the individual forecasting methods as the classes and the best method as the target class for each time series, enabling the use of existing classification algorithms. However, this simplification removes relevant information which may be useful in the metalearning process, such as which methods produce similar errors to the best one for a series, or which series are more *difficult* than others. We do not apply this simplification and take into account the exact error that each method produces in each series, instead of just considering which one produces the minimum error. This information is introduced into the model by posing the problem as finding a function that assigns *probabilities* to each forecasting method for each series, with the objective of minimizing the expected error that is produced if the methods were picked at random following these probabilities.

Our approach can be easily transformed into a classification exercise just by picking the method with the largest probability produced by our model. If we allow a sufficiently rich hypothesis space to fit freely, we also end up assigning probability one to the methods with the least amount of error. Our approach can also be seen as a classification exercise but with *per class weights* (the forecasting errors) that vary per instance, combined with *per instance weights* that assign more importance to some series.

2.2 Algorithmic description

The FFORMA framework has four main components:

1. The set of forecasting methods, the pool of methods.
2. The set of features to be extracted from the time series
3. A training set of time series, the reference set.
4. The forecasting error measure, such as mean squared error.

The FFORMA framework works in two phases: an **offline** phase, when the meta-learned is trained and an **online** phase, when forecasts are produced. Algorithm 1 describes the two phases.

Algorithm 1 The FFORMA framework - Forecast combination based on meta-learning.

Offline phase - train the learning model

Given:

$O = \{X_1, X_2, \dots, X_n\}$: the collection of N observed time series, the reference set.

P : Set of K forecasting algorithms such as ARIMA, ETS, SNAIVE, etc.

F : the set of functions to calculate time series features.

E : A forecasting error measure such as Mean Squared Error.

Output:

FFORMA meta-learner: A function from the extracted features to a set of K probabilities, one for each method in P .

Prepare the meta-data

For $j = 1$ to N :

- 1: Split X_j into a training period and test period.
- 2: Calculate the set of features for the training period by applying F .
- 3: Fit the models in P to the training period.
- 4: Calculate forecasts for the test period from each model.
- 5: Calculate forecast error measure E over the test period for all models in P .
- 6: Meta-data: input features x_j (step 2), output errors: e_j (step 5).

Train the meta-learner

- 7: Train a learning model based on the meta-data and errors, by minimizing:

$$\operatorname{argmin}_f \sum_{j=1}^N \sum_{k=1}^K f(x_j)_k e_{jk}$$

- 8: meta-learner.

Online phase - forecast a new time series

Given:

FFORMA classifier from step 8 .

Output:

Forecast for the new time series X_{new} .

- 9: For X_{new} calculate features x_{new} by applying F .
 - 10: From the features, use the meta-learner to produce w the vector of probabilities.
 - 11: Compute the individual forecasts of the methods in P for X_{new}
 - 12: Compute the final forecast by weighted average using w and the individual forecasts.
-

3 Implementation and Application to M4 Competition

3.1 Data preprocessing

For the M4 competition, we used the whole M4 dataset as the reference set. The M4 competition database consists of 100,000 real-world time series of yearly, quarterly, monthly, weekly, daily and hourly data. To divide the series into training and testing period, we used the forecast horizon specified in the database as the size of the temporal holdout. When the training period resulted too short (less than 2 periods), the size of the temporal holdout was reduced. Some series were removed because they were constant after removing the holdout part.

3.2 Time series features

We used a set of 42 features in for the model. The functions to calculate these features are implemented in `tsfeatures` R package by Hyndman et al. (2018). These features have been previously used by Talagala, Hyndman & Athanasopoulos (2018) and Hyndman, Wang & Laptev (2015). ?? provides a brief description of features used in this experiment. A detailed description of these features is provided in Talagala, Hyndman & Athanasopoulos (2018).

Question: calculation of features for time series with multiple seasonality, and short time series?

No exogenous features were used, even though they were available in the M4 dataset, such as which domain the series belongs to (e.g. macroeconomics, finance, tourism...).

3.3 Forecasting methods

We considered nine forecasting algorithms implemented in the `forecast` R package. They are (the specific R calls for fitting the models are given):

- i) automated ARIMA algorithm (`auto.arima`)
- ii) automated ETS algorithm (`ets`)
- iii) feed-forward neural network with a single hidden layer is fitted to the lags. The number of lags is automatically selected (`nnetar`)
- iv) random walk with drift (`rwf` with `drift=TRUE`)
- v) TBATS model (`tbats`)

Table 1: Features used for selecting a forecast-model.

	Feature	Description	Non-seasonal	Seasonal
1	T	length of time series	✓	✓
2	trend	strength of trend	✓	✓
3	seasonality	strength of seasonality	-	✓
4	linearity	linearity	✓	✓
5	curvature	curvature	✓	✓
6	spikiness	spikiness	✓	✓
7	e_acf1	first ACF value of remainder series	✓	✓
8	stability	stability	✓	✓
9	lumpiness	lumpiness	✓	✓
10	entropy	spectral entropy	✓	✓
11	hurst	Hurst exponent	✓	✓
12	nonlinearity	nonlinearity	✓	✓
13	alpha	ETS(A,A,N) $\hat{\alpha}$	✓	✓
14	beta	ETS(A,A,N) $\hat{\beta}$	✓	✓
15	hwalpha	ETS(A,A,A) $\hat{\alpha}$	-	✓
16	hwbeta	ETS(A,A,A) $\hat{\beta}$	-	✓
17	hwgamma	ETS(A,A,A) $\hat{\gamma}$	-	✓
18	ur_pp	test statistic based on Phillips-Perron test	✓	-
19	ur_kpss	test statistic based on KPSS test	✓	-
20	y_acf1	first ACF value of the original series	✓	✓
21	diff1y_acf1	first ACF value of the differenced series	✓	✓
22	diff2y_acf1	first ACF value of the twice-differenced series	✓	✓
23	y_acf5	sum of squares of first 5 ACF values of original series	✓	✓
24	diff1y_acf5	sum of squares of first 5 ACF values of differenced series	✓	✓
25	diff2y_acf5	sum of squares of first 5 ACF values of twice-differenced series	✓	✓
26	seas_acf1	autocorrelation coefficient at first seasonal lag	-	✓
27	sediff_acf1	first ACF value of seasonally-differenced series	-	✓
28	sediff_seacf1	ACF value at the first seasonal lag of seasonally-differenced series	-	✓
29	sediff_acf5	sum of squares of first 5 autocorrelation coefficients of seasonally-differenced series	-	✓
30	lmres_acf1	first ACF value of residual series of linear trend model	✓	-
31	y_pacf5	sum of squares of first 5 PACF values of original series	✓	✓
32	diff1y_pacf5	sum of squares of first 5 PACF values of differenced series	✓	✓
33	diff2y_pacf5	sum of squares of first 5 PACF values of twice-differenced series	✓	✓

- vi) Theta method forecast (thetaf)
- vii) naive forecasts (naive)
- viii) STLM-AR Seasonal and Trend decomposition using Loess with AR modeling of the seasonally adjusted series(stlm with modelfunction ar)
- ix) seasonal naive forecasts (snaive).

It is worthy to emphasize that we used the default parameters in most methods without any hand tuning. Only in `auto.arima` we specified a more thorough search for hyperparameters than its default version. In the case of an error when fitting the series (e.g. a series is constant), the SNAIVE forecast method is used instead.

Question: i) Calculation of ets models to daily, hourly and weekly series

3.4 Metal-learning model

We chose the gradient tree boosting model of **xgboost** (Chen & Guestrin (2016)) as the underlying implementation of the learning model for the following reasons:

1. Ability to customize the model to fit our objective function
2. Computational efficiency: The size of the M4 dataset requires a efficient model
3. Good performance in structure based problems

3.4.1 The xgboost objective function

The vanilla xgboost algorithm produces numeric values from the features, one for each forecasting method in our pool. We apply the softmax transform to these values prior to computing the the objective function shown in Algorithm 1 step 7, implemented as a custom objective function.

xgboost fits the model by gradient boosting, for which requires a gradient and hessian of the objective function. In our case, the gradient is the direct gradient of the objective, but the *correct* hessian is prone to numerical problems that need to be fixed for xgboost to converge. This is a relative common problem and one simple fix is to use an upper bound of the hessian by clamping small values to a larger one. In our case, we compute an alternate upper bound of the hessian by removing some terms from the *correct* hessian. Although both alternatives converge, our approach works faster, requiring less boosting steps to converge. This not only increases the computational efficiency, it also generalizes better due to a less complex set of trees produced in the final solution.

The functions involved in computing the objective are the following:

- $y(x)$ is the output of the xgboost algorithm
- $p_j = \frac{e^{y(x)_j}}{\sum_k e^{y(x)_k}}$ is the transformation to probabilities by applying the softmax transform.
- $L = \sum p_j e_j$ is the loss of the function p .
- $G_j = \frac{\partial L}{\partial p_j} = p_j(e_j - \sum_k e_k p_k)$

$$H_j = \frac{\partial G_j}{\partial p_j} \approx \hat{H}_j = p_j(e_j(1 - p_j) - G_j)$$

3.4.2 Hyperparameters

The results of xgboost are particularly dependent on its hyperparameters such as learning rate, number of boosting steps, maximum complexity allowed for the trees or subsampling sizes. In our case we limit the hyperparameter search space based on some initial results and rules of thumb and explore it using bayesian optimization, (implemented in the `rBayesianOptimization` R package) measuring performance on a 10% holdout version of the reference set. We picked the most simple hyperparameter set from the top solutions of the exploration.

3.5 Prediction Intervals

The M4 Competition also featured a subcompetition over the accuracy of prediction intervals. For this part, we used a different approach. In order to compute the intervals we used the point forecast produced by our meta-learner as the centre of the interval and computed the 95% bounds of the interval by a linear combination of the bounds of three forecasting methods: Thetaf, SNAIVE and NAIVE methods. The coefficients for the linear combination were calculated in a data driven way also over the M4 database. The procedure is as follows:

1. For the training period each series in the dataset:
 1. Compute the point forecast of the meta-learning.
 2. Compute the 95% *predicion radius* for the thetaf, snaive and naive, this is the difference between the 95% upper bound and the point forecast for each horizon.
2. For each forecasting horizon required in the dataset:
 1. Find the coefficients that minimize the MSIS error of the interval with the meta-learning point forecast as center and a linear combination of the radiuses of thetaf, snaive and naive as radius. The minimization is done by gradient descent.

This procedure produces at set of three coefficients for each prediction horizon in the M4 dataset and these coefficients will be the same independently of the series we want to forecast. *These coefficients are not restricted to be probabilities, the optimization is unrestricted.* In order to prevent overfitting for the center of the intervals, the M4 dataset was divided in two halves, and our metalearning approach was trained in one half and applied to the other.

4 Results

We quantify the improvement in accuracy produced by our approach, compared to a classification implementation using the same underlying implementation, xgboost.

References

- Bates, JM & CWJ Granger (1969). The Combination of Forecasts. *OR* **20**(4), 451–468.
- Chen, T & C Guestrin (2016). Xgboost: A scalable tree boosting system. In: *Proceedings of the 22nd acm sigkdd international conference on knowledge discovery and data mining*. ACM, pp.785–794.
- Hyndman, RJ, E Wang & N Laptev (2015). Large-scale unusual time series detection. In: *Data Mining Workshop (ICDMW), 2015 IEEE International Conference on*. IEEE, pp.1616–1619.
- Hyndman, R, E Wang, Y Kang & T Talagala (2018). *tsfeatures: Time Series Feature Extraction*. R package version 0.1. <https://github.com/robjhyndman/tsfeatures>.
- Talagala, TS, RJ Hyndman & G Athanasopoulos (2018). Meta-learning how to forecast time series. *Technical Report 6/18, Monash University*.
- Timmermann, A (2006). “Forecast Combinations”. In: *Handbook of economic forecasting*, pp. 135-196. Chapter 4. Ed. by Graham Elliot and Clive W. J. Granger, Timmermann Allan. Amsterdam, North-Holland.